

# Creating and Using the OpenStack Aware Network

**ARISTA**

# Agenda

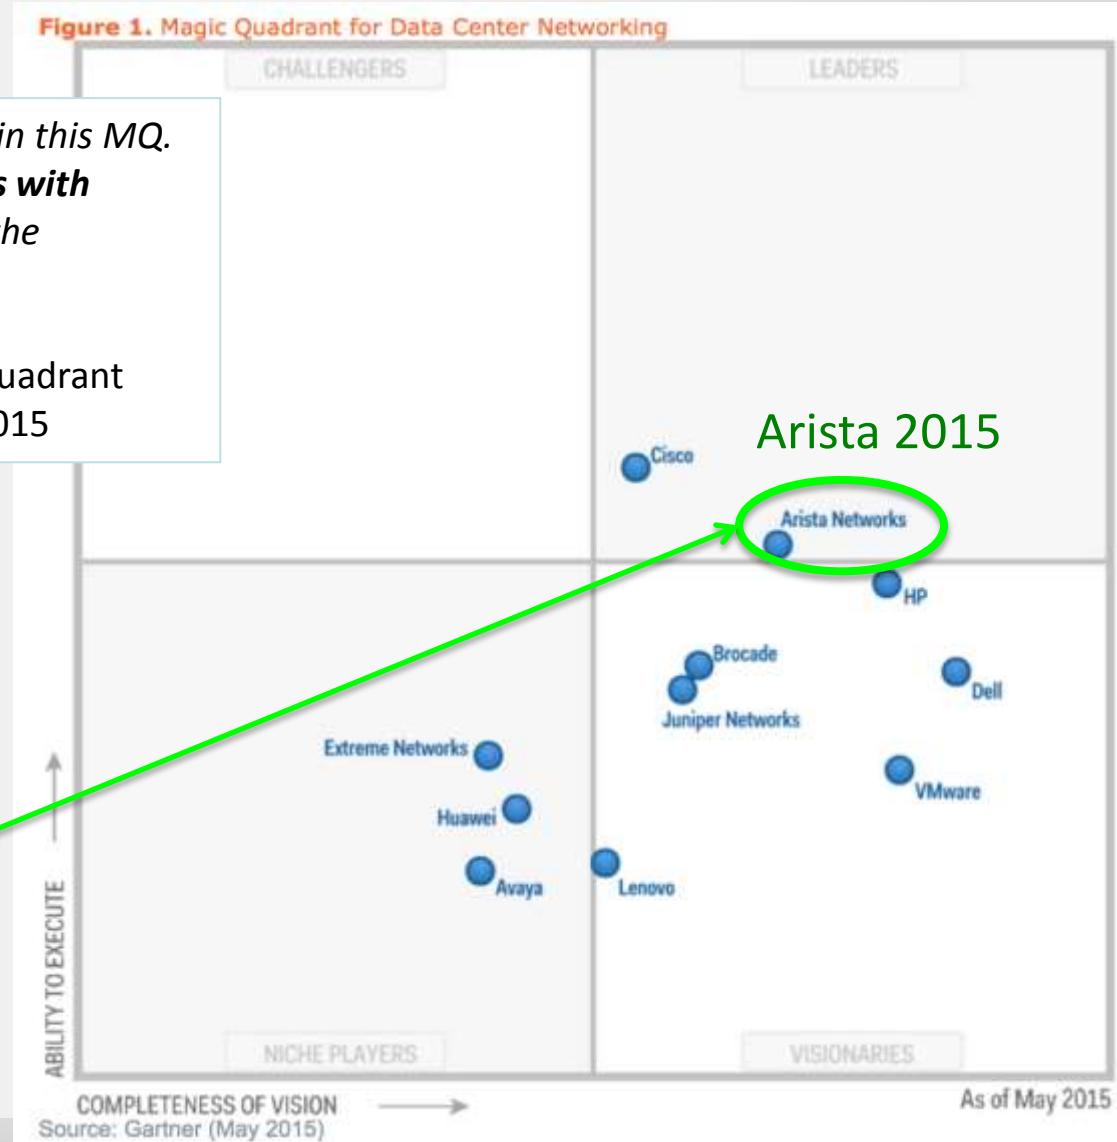
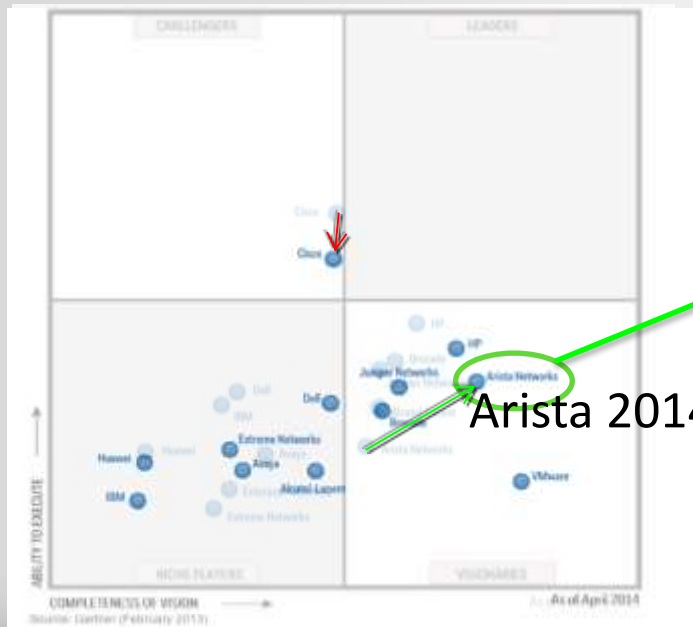
- Company overview
- Drive for Network Virtualization
- Components of Network Virtualization stack
- Concept of CVX
- OpenStack
- OpenStack Demo

# 2015 Gartner MQ Data Center Networking


Arista placed in the **leadership** quadrant

1. Arista is by far the **fastest-growing** vendor in this MQ.
2. Arista provides **high-performance solutions with deep buffers and low latency** to deal with the complexities of modern DC applications.

Gartner Data Center Networking Magic Quadrant  
May 11, 2015



# About Arista Networks



## 10/40/100GbE Networks for the Virtualized Cloud & Data Center

- Founded in 2004
- Shipping Since Mid-2008
- NYSE: ANET in 2014/6
- 3000+ Customers
- 1000+ Employees

Profitable, self-funded network infrastructure provider

Founded to build the best Network Operating System for Next Generation Data Centers



# Key Executives



Andy Bechtolsheim

Founder, Chief Development Officer and Chairman

Founder of Sun Microsystems

Founder of Granite Systems

Initial investor in Google, Inc.



Jayshree Ullal

President and CEO

Part of Cisco's 1<sup>st</sup> acquisition of Crescendo

15 yr Cisco SVP for Data Center, Switching & Services

Joined Arista in 2008

# Arista : The Best Data Center Portfolio

Extensible Operating System(EOS) - Single image

Openstack Integration for all platform



## 7010T & 7048T

48-port Data Center  
Class Gigabit  
Ethernet Switch

## 7150S

Ultra Low Latency  
24,52,64-port SFP+  
1G-40GbE Switches  
LANZ and DANZ

## 7050X & 7250X

Dense Low Latency  
32 & 64-port QSFP

96xSFP+/8xQSFP  
Advanced  
Virtualization  
Scale-out  
Visibility

## 100G

### 7280SE

10/40/100G

Ultra Deep Buffers  
VOQ and Lossless  
Enhanced Visibility  
LANZ/DANZ  
NEBS

### 7300X

High Density,  
Modular System  
supporting up to 512  
40GbE

Cloud Scale  
Leaf and Spine  
10/40G

## 100G

### 7500E


Lossless, High  
Density, Modular  
Switching System  
supporting up to  
1152 Wire speed  
10GbE Ports  
LANZ / DANZ  
Spine  
10/40/100G

# ARISTA



# Benefits of EOS Architecture

Other Network OS



VS

Arista EOS



Linux Kernel

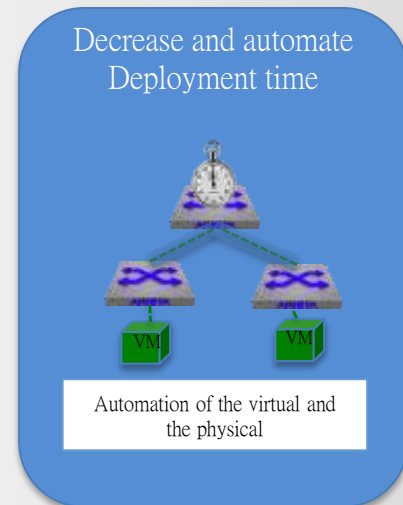

Data Plane

ARISTA

# Network Virtualization

## ■ Goals and drive for Network Virtualization

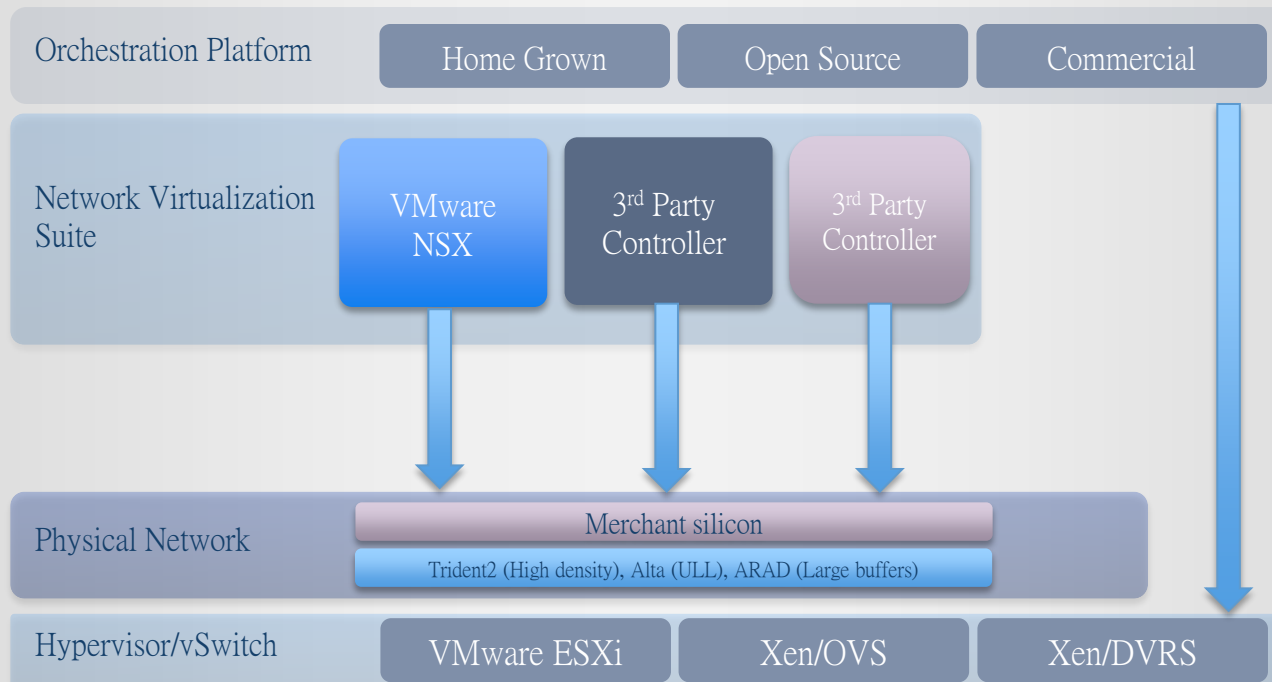
- Remove underutilized islands of resources
- Single resource pool to allow optimal utilization of all available resources
- Provide any-to-any L2 or L3 communication within the Data Center
- Deployment and connect resources in seconds rather than days/weeks





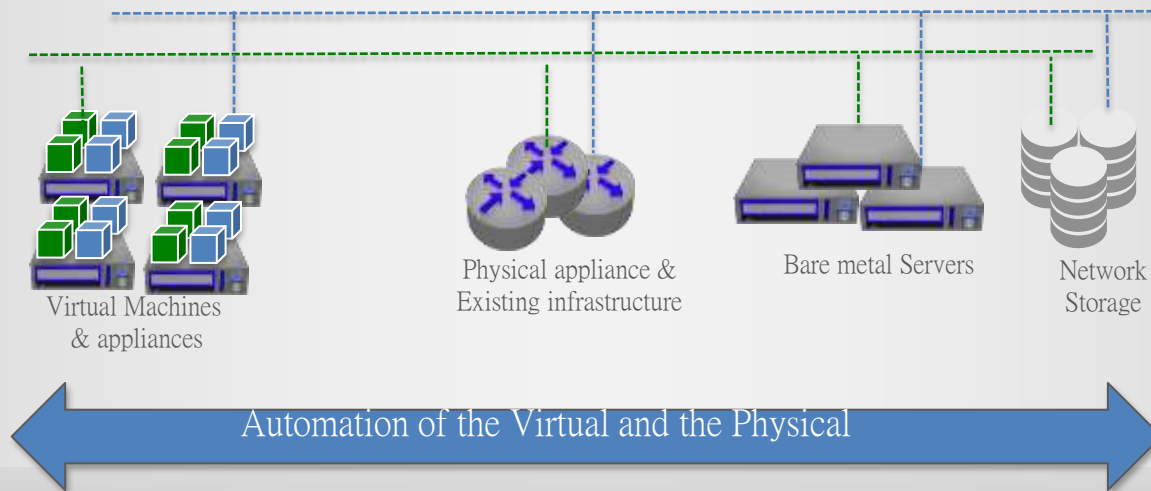
# Network Virtualization

- Components of a Network Virtualization stack



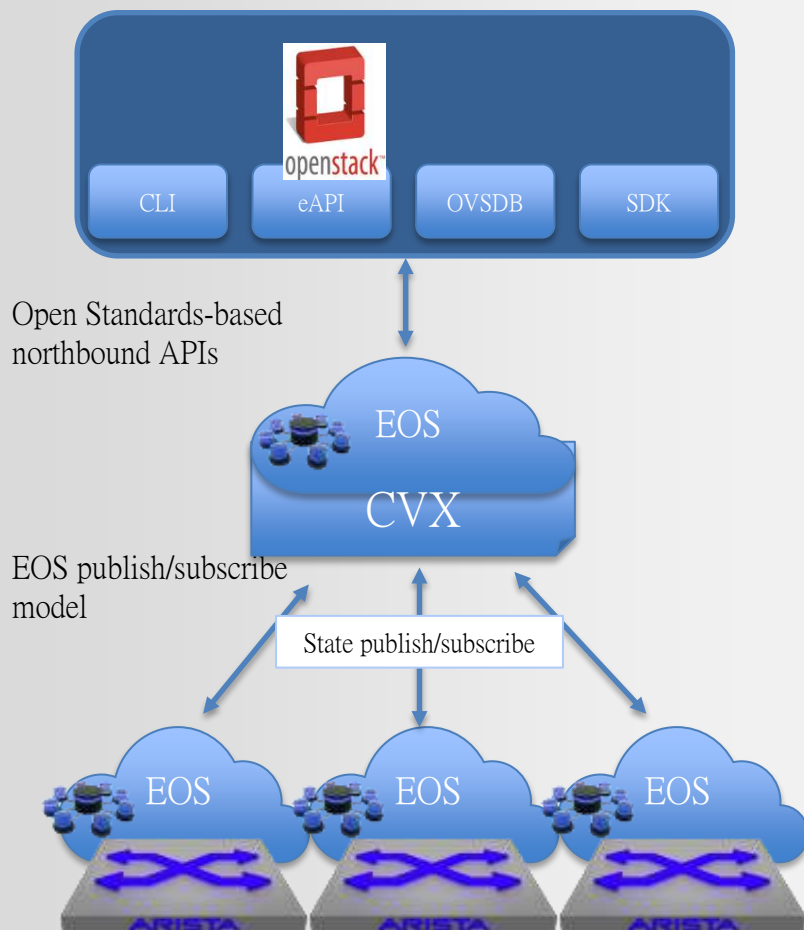
# Network Virtualization

- Requirement to automate connectivity for Physical appliances
  - Virtualization stack today or in the future, will need to provide connectivity to the physical infrastructure
  - Provide connectivity for Bare Metal Server, Storage, DC routers
  - Connectivity needs to be provided in an automated manner, similar to the service delivered for virtual appliances



# Introducing CloudVision eXchange (CVX)

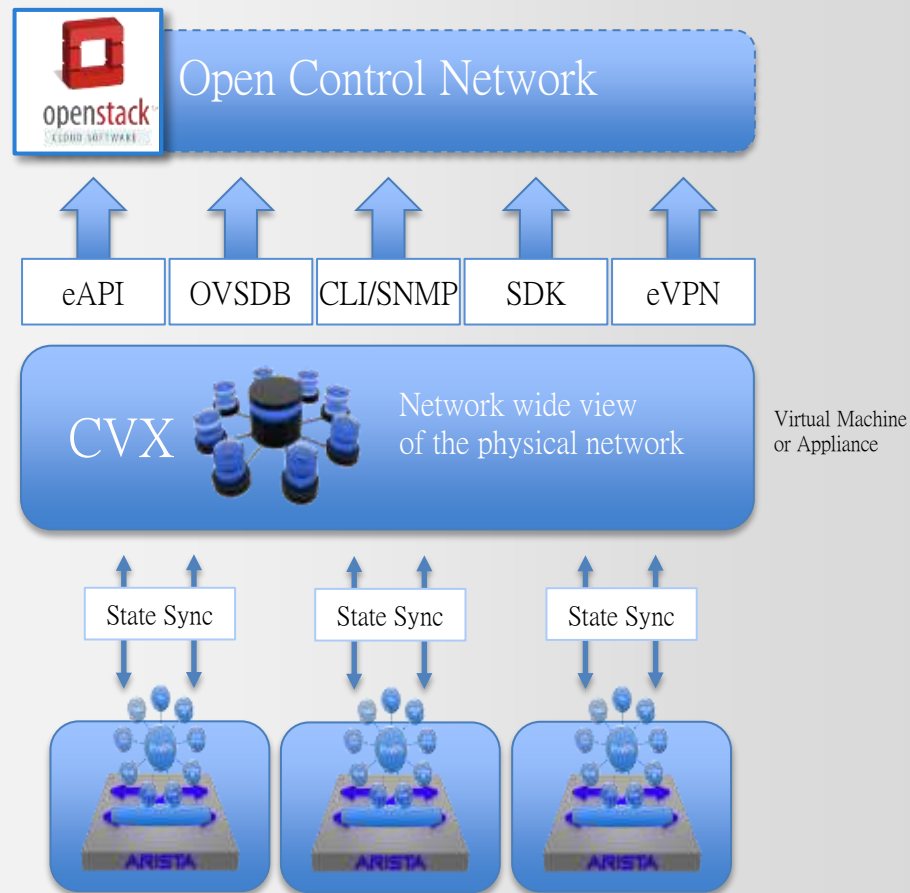
Single Interface to EOS devices




- Part of EOS CloudVision Framework
- Abstracts the physical switch infrastructure
- Provides a single access point for real-time provisioning, orchestration and integration with multi-vendor controllers
- Distributed EOS state: CVX mounts state from all switches (sysDBs) in the network
- No new protocol needed – uses EOS framework
- Management plane, not data plane
- May be a VM or Appliance

# CVX – more detail

- CVX is software Agent enabled on the EOS operating System.
- CVX mounts and synchronizes configurable state from the physical switches
- Providing CVX with an aggregated view of the network state
- The state mounted is depended on the service being deployed
- Each service specifies what state it mounts from / pushes to each switch
- Services include:
  - VCS (VXLAN Control Service)
  - OpenStack
  - Network Topology

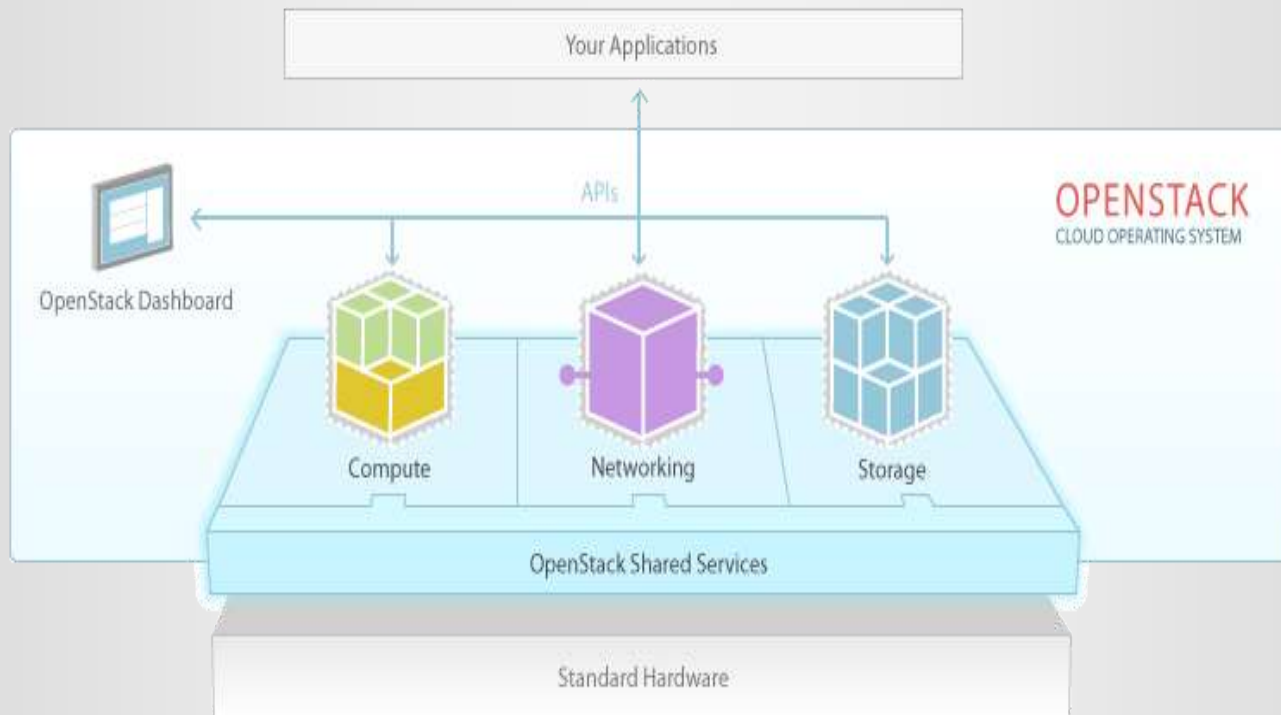




# OpenStack Overview

**ARISTA**

# OpenStack Overview



OpenStack is a **cloud operating system** that controls large pools of compute, storage, and networking resources throughout a datacenter, all managed through a dashboard that gives administrators control while empowering their users to provision resources through a web interface.



# OpenStack Core Services



- Compute ("Nova") provides virtual servers upon demand.



- Network ("Neutron") is a pluggable, scalable and API-driven system for managing networks and IP addresses.




- Block Storage ("Cinder") provides persistent block storage to guest VMs

# OpenStack Core Services



- Dashboard ("Horizon") provides a modular web-based user interface for all the OpenStack services



- Object Store ("Swift") provides object storage. It allows you to store or retrieve files
- Image ("Glance") provides a catalog and repository for virtual disk images
- Identity ("Keystone") provides authentication and authorization for all the OpenStack services



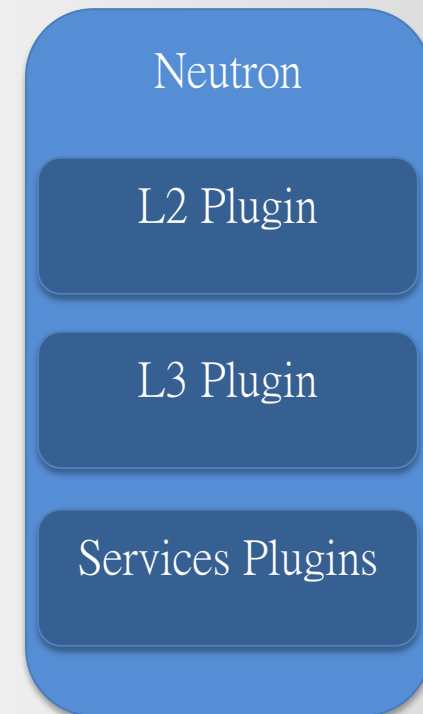
# OpenStack Development Cycle

- From the Diablo release onwards, 6-month release cycle
  - Two releases a year, ~April and October
  - Release names are alphabetic and based on the cities/counties where the corresponding OpenStack summit took place

Release	Release Date	Status
Folsom	September 2012	EOL
Grizzly	April 2013	EOL
Havana	October 2013	EOL
Icehouse	April 2014	Security Supported (patches built)
Juno	October 2014	Previous release
Kilo	April 2015	Current release
Liberty	Under Discussion	Q4 2015

# Neutron Overview


- OpenStack Networking Service
  - Introduced in the Folsom release
- Manages
  - Tenant Networks
  - Logical Routers
  - L4-7 Services – NAT, LBAAS, FWAAS, ...



# Modular Layer 2 (ML2) Plugin

- Neutron plugin introduced in Havana which provides:
  - Separation between the state of tenant networks and how that state is then realized across the network
  - Flexibility in how the virtual and physical network are managed
  - Multi-vendor support via multiple “Mechanism Drivers” managing pieces of the network in parallel
- “Monolithic” plugins are being deprecated in favor of ML2





# Arista CVX Integration with OpenStack

**ARISTA**

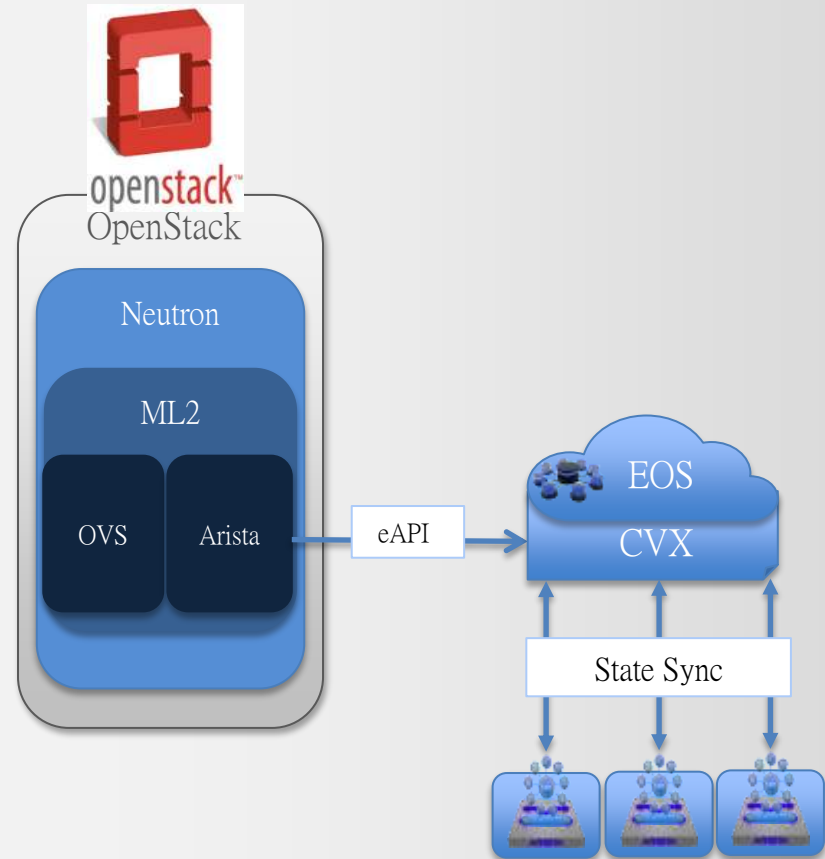


# Arista OpenStack Integration

- Arista Focus - Provide end-to-end tenant network visibility and provisioning across virtual and physical infrastructure for the highest performance VLAN and VXLAN fabrics
- Key Solution Features
  - Provide full visibility to the network team of the OpenStack state
  - Automatically provision the physical network in response to tenant configuration within OpenStack

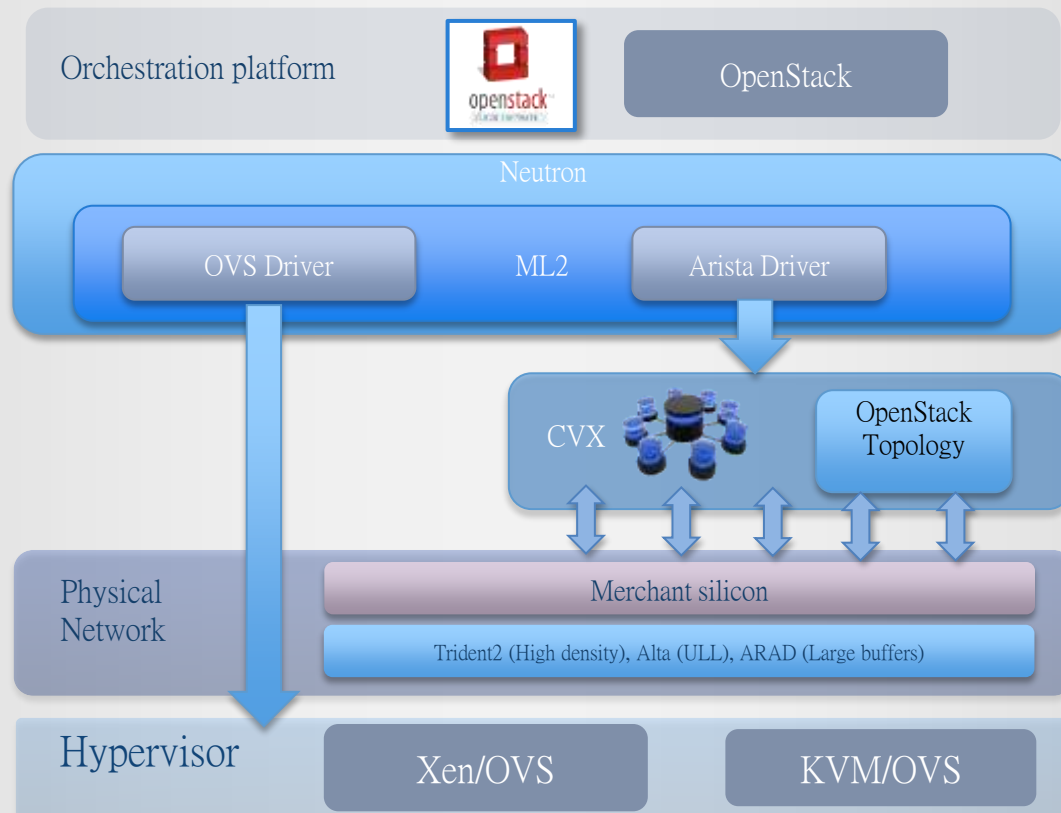
# Arista OpenStack Integration

- Arista CVX
  - Collection of global network topology from switches
  - Provisioning of the switches via eAPI
- Arista ML2 mechanism driver
  - Notify created OpenStack state to the CVX
  - Communication via the Arista eAPI



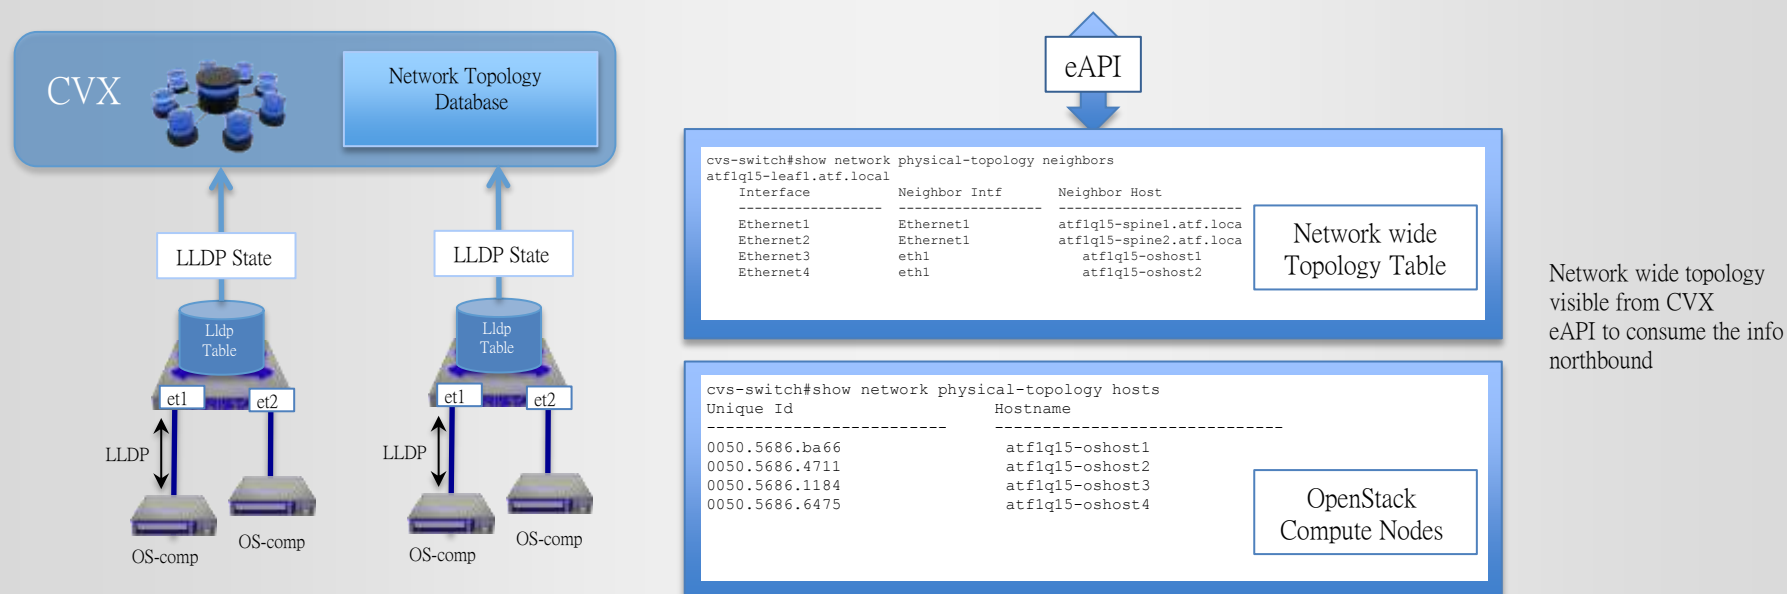
# Arista OpenStack Integration

- Network orchestration with OpenStack Neutron ML2 Plug



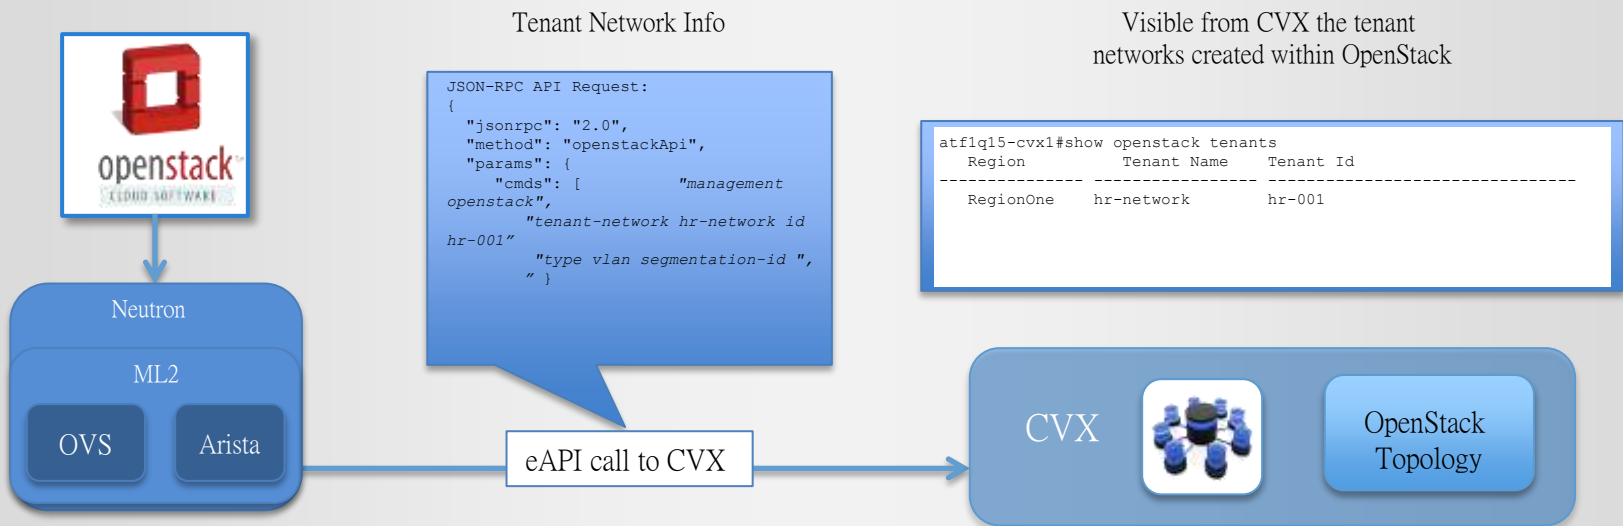
# Arista OpenStack Integration

- CVX builds a full topology of the physical infrastructure
  - Leaf switch builds their local topology table using standard LLDP
  - Contains directly attached OpenStack compute nodes, which will host the virtual machines
  - CVX mounts the local LLDP tables, providing a network wide view
  - CVX knows the physical location (switch and interface) each OS compute node is attached



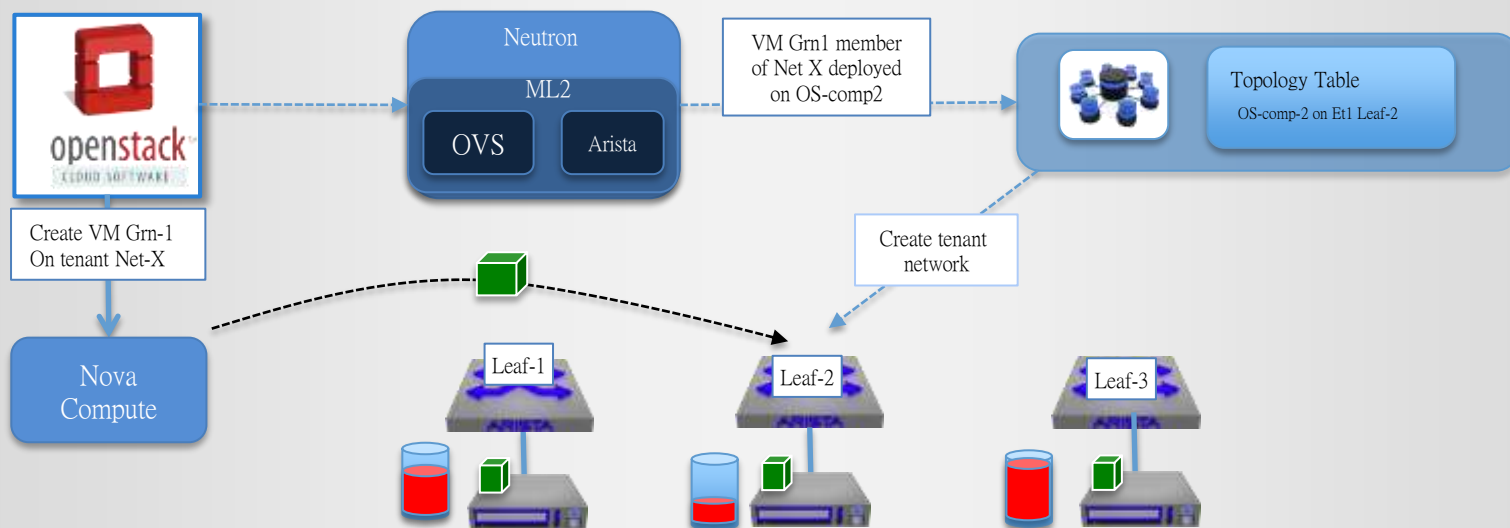
# Arista OpenStack Integration

- Step1: OpenStack, Horizon portal create a tenant network
  - ML2 plugin with the Arista driver installed
  - Plugin converts the network create call to an eAPI to CVX
  - CVX is now aware of all networks created with OpenStack
  - Network only deployed on the physical infrastructure if a VM is attached to the network



# Arista OpenStack Integration


- Step 2: OpenStack, Horizon portal create VM and apply to network
  - Virtual Machine created within Horizon and attached to the tenant's network
  - Nova-scheduler deploys the VM on the compute node based on available resources
  - CVX has visibility of the compute node the VM is deployed on via the plug
  - From the global LLDP table maps the Compute node to a switch port, and dynamically deploys the tenant network on the switch





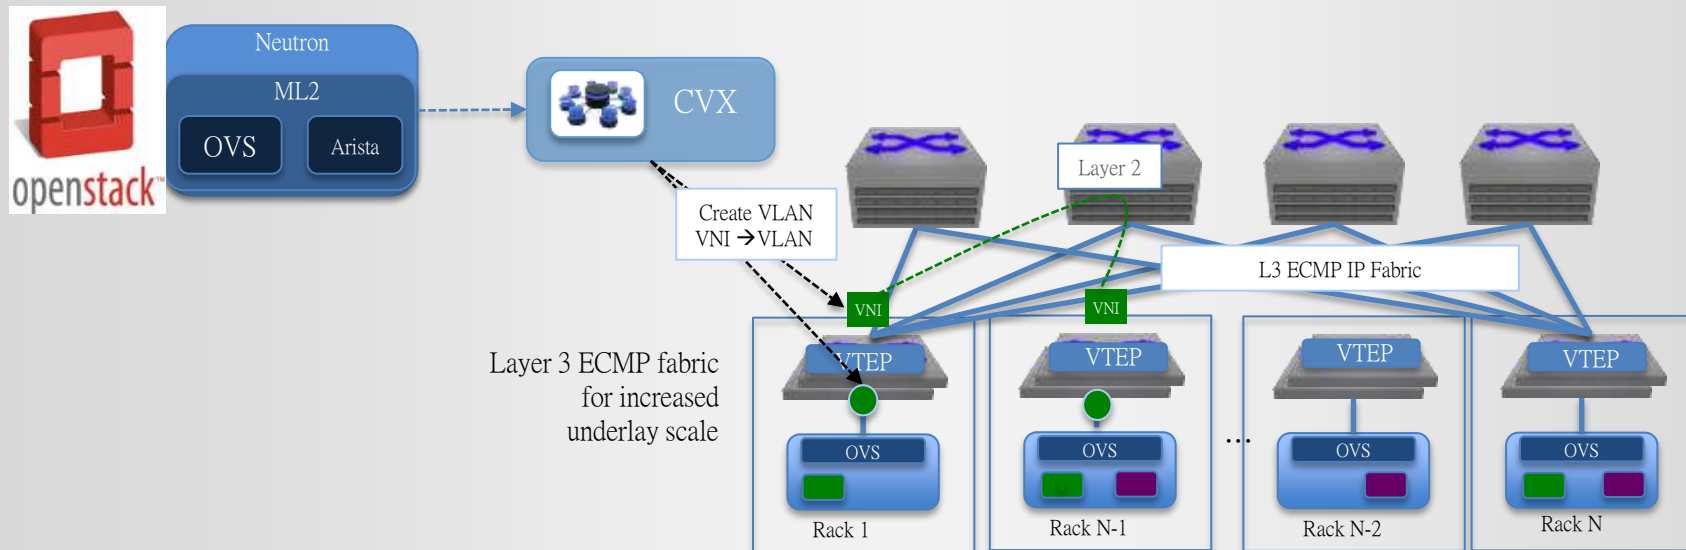
# Arista OpenStack Integration – VLAN Based

- Network Type created by the CVX can be a simple Layer 2 VLAN
  - Tenants VMs deployed within dedicated tenant VLANs
  - CVX dynamically deploys the tenant VLAN on the appropriate Leaf edge port and Uplink
  - Deployed on a standard MLAG (Layer 2) topology



# Arista OpenStack Integration – VXLAN Based

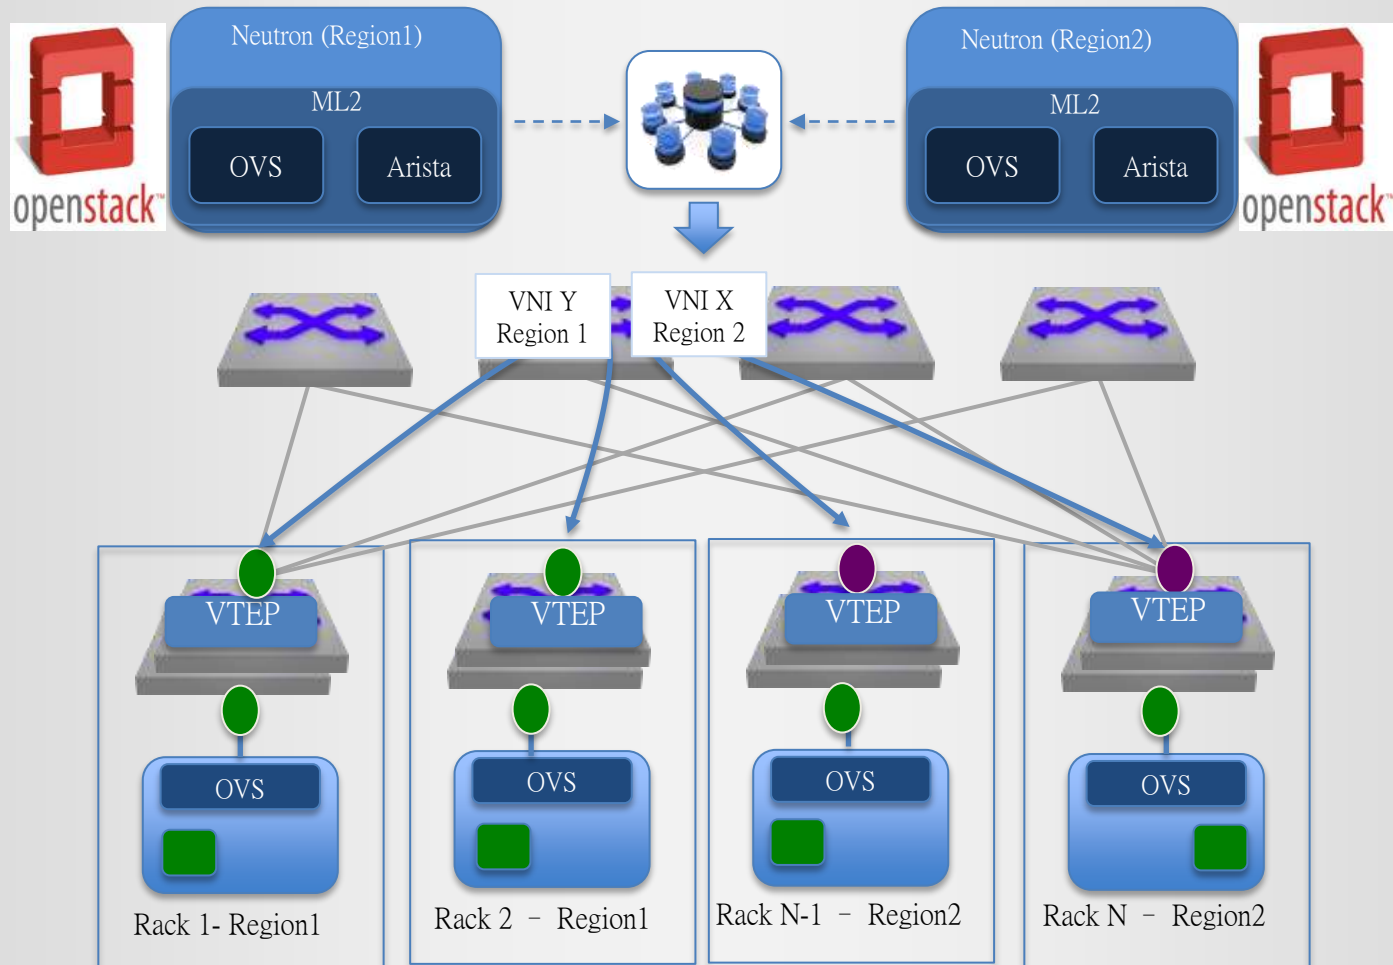
- OpenStack integration with VXLAN
  - Scalable IP fabric with a Layer 3 ECMP design
  - Hardware VXLAN VTEP configured on every leaf switch
  - Layer 2 connectivity between rack via VXLAN across the L3 fabric



# Scaling OpenStack

- Multiple OpenStack clusters supported per CVX instance
- Can be combined with other network virtualization
  - NSX
  - Etc
- VXLAN breaks out of the 4K VLAN limit
  - 16M VNIs mapped to locally significant VLANs

# Multi-Tenant OpenStack Deployment




# Routing with OpenStack

- L2 up until now, how do you route?
- Can be performed by a Network Node
- Allows connectivity between tenants and external networks
  - NAT Support
  - VRF Support
- Limited by software
- Alternative is perform this at the switch...with limitations!

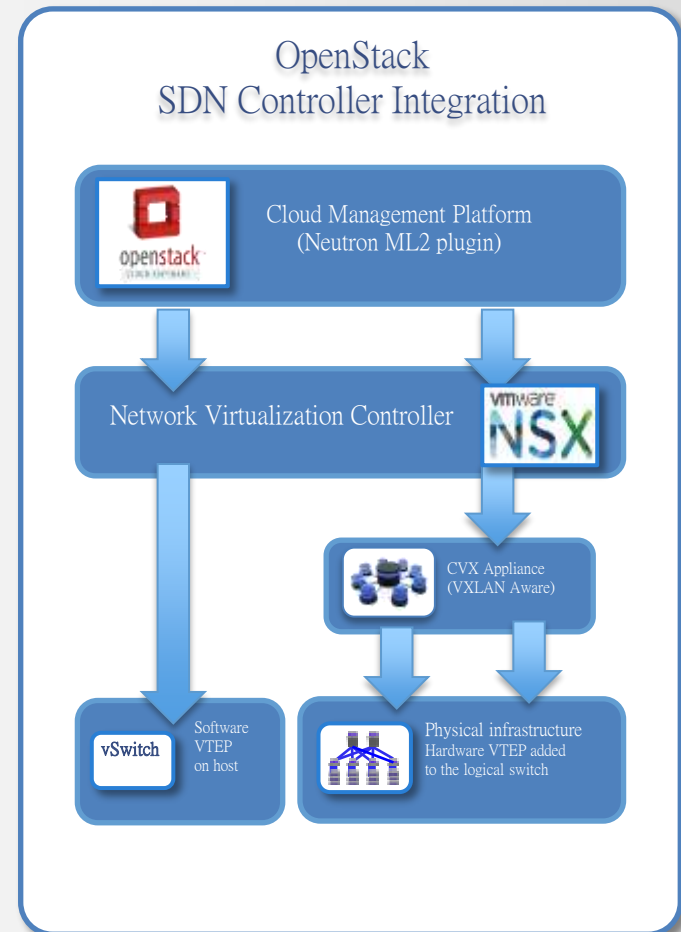
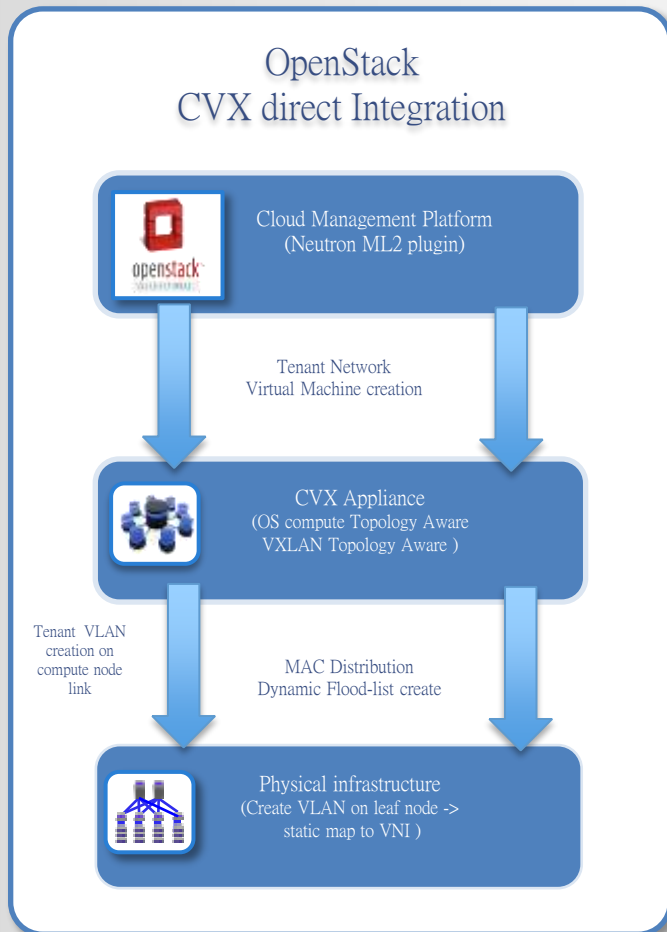
# OpenStack Integration – L3 Plugin

- Arista L3 plugin provisions SVIs over eAPI in response to tenant's creating logical routers
- Routing happens at dedicated network nodes
  - Pair of MLAGed physical devices
  - Active-Active HA via MLAG
  - Performs routing for the OpenStack cluster
    - Can be scaled out horizontally by tenant as needed
- TORs can also be used as the routing nodes

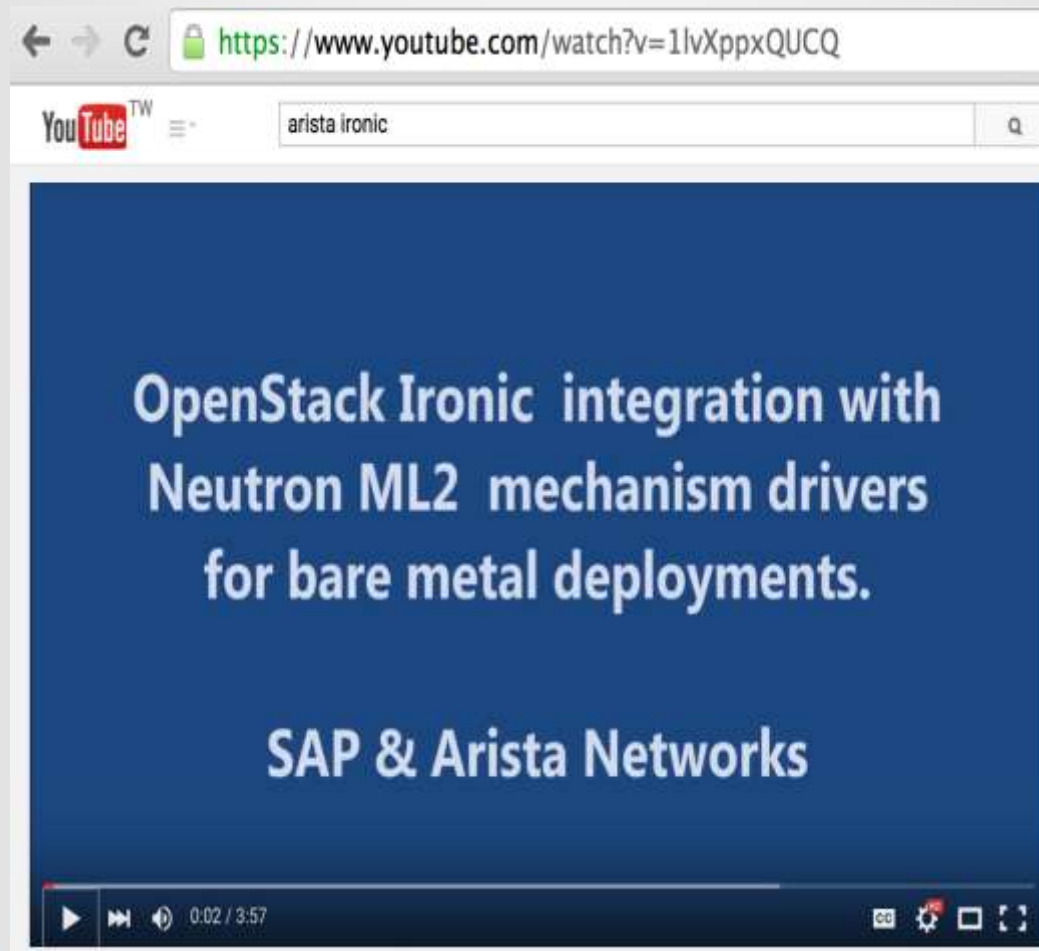




# OpenStack Orchestration models




# Openstack Ironi Integration



<https://www.youtube.com/watch?v=1lvXppxQUCQ>

ARISTA




OpenStack Demo

**ARISTA**

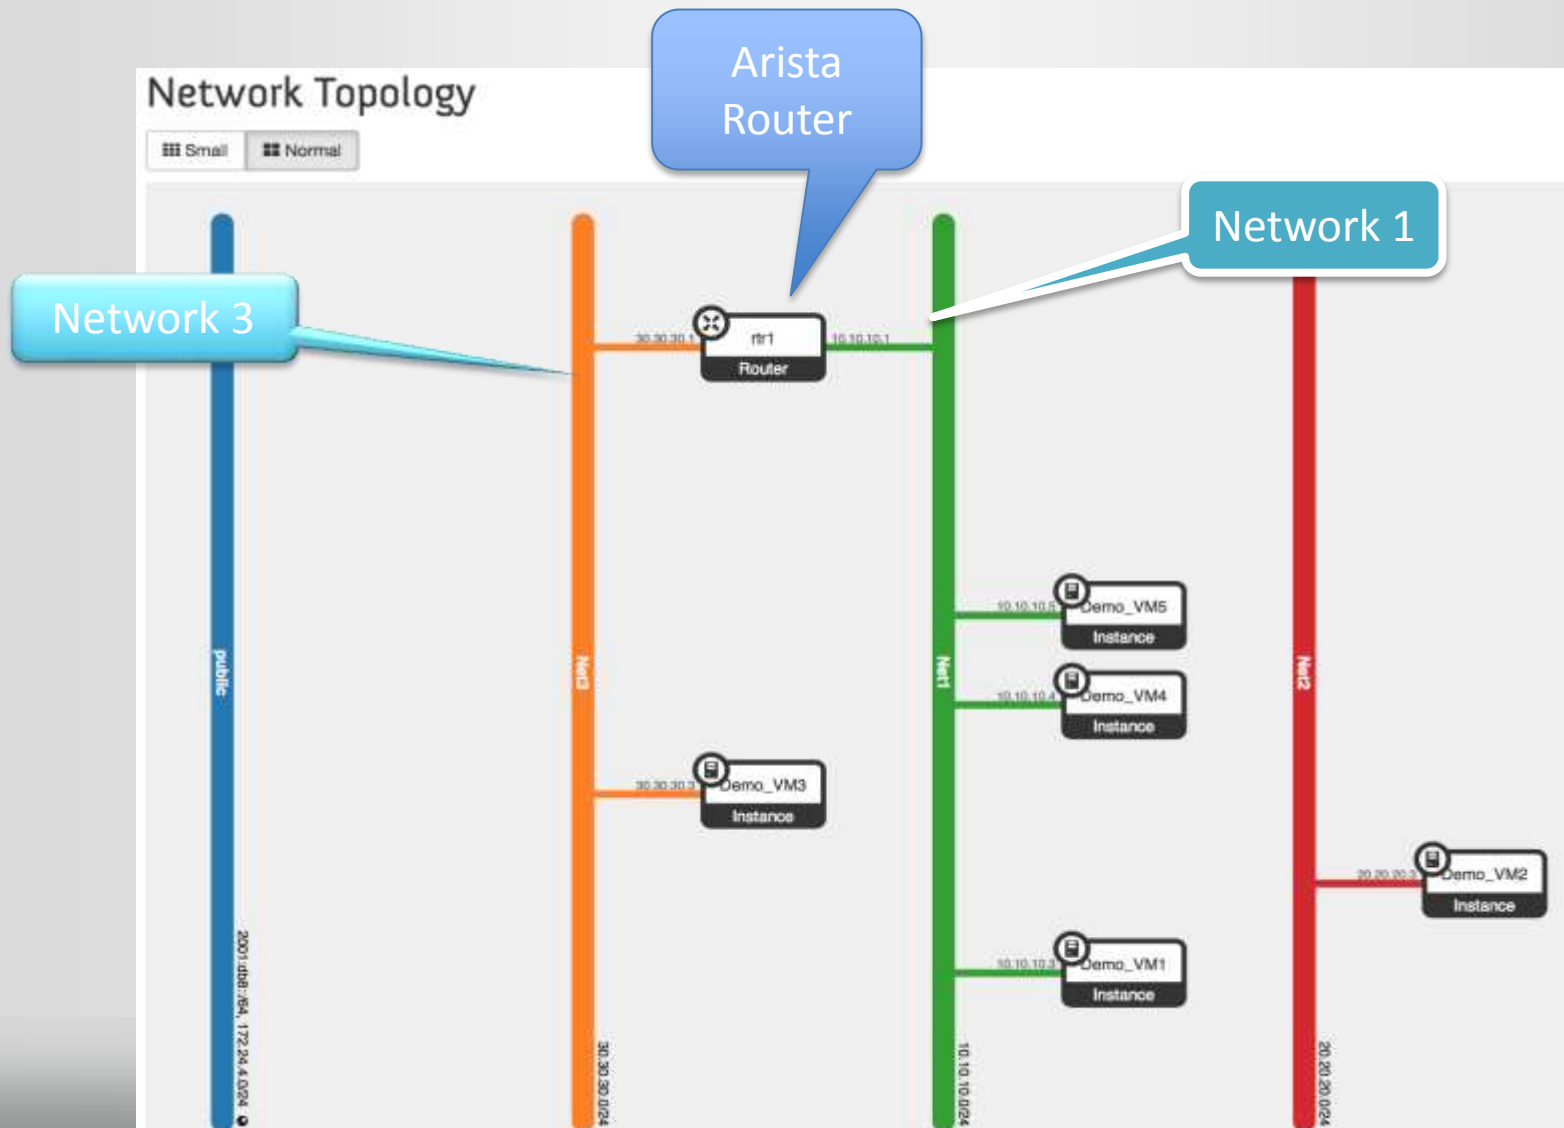
# Arista OpenStack Integration – VXLAN Based Network Architecture

## Demo Diagram



# Openstack Dashboard

## Network Topology 1



# Create Network and Instance

## Network Details: Net1

### Network Overview

Name	Net1
ID	8b950923-6b52-4e16-8257-49c4c8f294c1
Project ID	c832cc8c4739426cbb0c08824fabe10b
Status	ACTIVE
Admin State	UP
Shared	No
External Network	No
MTU	Unknown
Provider Network	Network Type: vlan Physical Network: default Segmentation ID: 1174



# Create Network

The screenshot shows the OpenStack dashboard interface. The sidebar on the left contains navigation links: Project, Compute, Network, Network Topology, Networks (highlighted in red), Routers, Admin, and Identity. The main content area is titled 'Networks' and features a search filter and a '+ Create Network' button. Below this is a table listing the existing networks.

	Name	Subnets Associated	Shared	Status	Admin State
<input type="checkbox"/>	Net3	30.30.30.0/24	No	Active	UP
<input type="checkbox"/>	public	ipv6-public-subnet 2001:db8::/64 public-subnet 172.24.4.0/24	No	Active	UP
<input type="checkbox"/>	Net1	10.10.10.0/24	No	Active	UP
<input type="checkbox"/>	Net2	20.20.20.0/24	No	Active	UP
<input type="checkbox"/>	Medialinks_net1	192.168.155.0/24	No	Active	UP

# Create Network

## Create Network

Network Subnet\* Subnet Details

Network Name

Create a new network. In addition, a subnet associated with the network can be created in the next panel.

Admin State ?

UP

Cancel

« Back

Next »

## Create Network

Network Subnet Subnet Details

☒ Create Subnet

Subnet Name

Network Address \* ?

IP Version \*

IPv4

Gateway IP ?

☐ Disable Gateway

Create a subnet associated with the new network, in which case "Network Address" must be specified. If you wish to create a network without a subnet, uncheck the "Create Subnet" checkbox.

Cancel

« Back

Next »



# Launch Instance

openstack admin ▾

## Instances

Instance Name ▾ Filter Filter [Launch Instance](#)

<input type="checkbox"/>	Instance Name	Image Name	IP Address	Size	Key Pair	Status	Availability Zone	Task	Power State
<input type="checkbox"/>	<a href="#">Medialinks_vm1</a>	cirros-0.3.4-x86_64-uec	192.168.155.3	<a href="#">m1.nano</a>	-	Active	nova	None	Running
<input type="checkbox"/>	<a href="#">Demo_VM5</a>	cirros-0.3.4-x86_64-uec	10.10.10.5	<a href="#">m1.nano</a>	-	Active	nova	None	Running
<input type="checkbox"/>	<a href="#">Demo_VM4</a>	cirros-0.3.4-x86_64-uec	10.10.10.4	<a href="#">m1.nano</a>	-	Active	nova	None	Running
<input type="checkbox"/>	<a href="#">Demo_VM3</a>	cirros-0.3.4-x86_64-uec	30.30.30.3	<a href="#">m1.nano</a>	-	Active	nova	None	Running
<input type="checkbox"/>	<a href="#">Demo_VM2</a>	cirros-0.3.4-x86_64-uec	20.20.20.3	<a href="#">m1.nano</a>	-	Active	nova	None	Running
<input type="checkbox"/>	<a href="#">Demo_VM1</a>	cirros-0.3.4-x86_64-uec	10.10.10.3	<a href="#">m1.nano</a>	-	Active	nova	None	Running

# Launch Instance

Launch Instance

Details \*

Access & Security

Networking \*

Post-Creation

Advanced Options

Availability Zone

nova

Instance Name \*

Demo\_VM1

Flavor \* ?

m1.nano

Instance Count \* ?

1

Instance Boot Source \* ?

Boot from image

Image Name \*

cirros-0.3.4-x86\_64-uec (24.0 MB)

Specify the details for launching an instance.

The chart below shows the resources used by this project in relation to the project's quotas.

Flavor Details

Name	m1.nano
VCPUs	1
Root Disk	0 GB
Ephemeral Disk	0 GB
Total Disk	0 GB
RAM	64 MB

Project Limits

Number of Instances

6 of 10 Used

Number of VCPUs

6 of 20 Used

Total RAM

334 of 51,200 MB Used

Cancel

Launch

## Launch Instance

[Details \\*](#) [Access & Security](#) [Networking \\*](#) [Post-Creation](#) [Advanced Options](#)

### Selected networks

Choose network from Available networks to Selected networks by push button or drag and drop, you may change NIC order by drag and drop as well.

NIC:1 Net1

# Display Network information on Arista switch

```
bleaf1#show openstack networks
```

```
Region: RegionOne
```

```
Tenant Name: admin
```

```
Tenant Id: c832cc8c4739426cbb0c08824fabe10b
```

Network Name	Network Id	Seg Type	Seg Id	Maps to VNI
Medialinks_net1	ed297f0d-ea00-474d-8726-17d2c7b10dd6	vlan	1145	11145
Net1	8b950923-6b52-4e16-8257-49c4c8f294c1	vlan	1174	11174
Net2	963448c9-a685-4a73-ad76-0e63aba59463	vlan	1190	11190
Net3	131b0b55-f4dc-4602-9510-6eb744a1bd3a	vlan	1143	11143
public	4f2810e6-9940-4308-b5c9-8c82a86335bb	vlan	1194	11194

# Display VM visibility on Arista switch (1)

```
bleaf1#show openstack vms
```

```
Region: RegionOne
```

```
Tenant Name: admin
```

```
Tenant Id: c832cc8c4739426cbb0c08824fabe10b
```

VM Name	VM Id	Host	Network Name
Demo_VM1	12fc3936-29c4-47a4-b8b3-33b73b04d75c	kilo1	Net1
Demo_VM2	6ad8e5d3-41ae-41c3-a449-48681173ee1f	kilo2	Net2
Demo_VM3	14ecafb6-6011-48f7-ab3f-9ffdf3475ce8	kilo1	Net3
Demo_VM4	6bc4eeb3-3ed9-43da-9c6f-b8b3b6885fc3	kilo2	Net1
Demo_VM5	f2c4f755-df2b-49b0-9e47-3a221ad759e6	kilo1	Net1

## Display VM visibility on Arista switch (2)

```
bleaf1#show openstack vms vm Demo_VM1 detail  
Region: RegionOne
```

```
Tenant Name: admin  
Tenant Id: c832cc8c4739426cbb0c08824fabe10b
```

```
VM Name: Demo_VM1
```

```
VM Id: 12fc3936-29c4-47a4-b8b3-33b73b04d75c
```

```
Host: kilo1
```

Switch Name	Switch Id	Switch Interface	Seg Type	Seg Id
bleaf1.aristanetworks.com	001c.7300.4406	Ethernet18	vlan	1174

# Display physical-network visibility on Arista switch

```
bleaf1#show openstack physical-network switch bleaf1.aristanetworks.com
```

```
Switch: 001c.7300.4406 ( bleaf1.aristanetworks.com )
```

```
Ethernet18:
```

```
Allowed vlans: 1120,1143,1145,1174,1190
```

```
Connected host: kilo1
```

```
Region: RegionOne
```

```
Tenant: admin ( c832cc8c4739426cbb0c08824fabe10b )
```

```
...skipping...
```

```
VM: Demo_VM1 ( 12fc3936-29c4-47a4-b8b3-33b73b04d75c )
```

Network Name	Network Id	Seg Type	Seg Id
Net1	8b950923-6b52-4e16-8257-49c4c8f294c1	vlan	1174



# VXLAN enable con Arista switch

```
bleaf1#show int vxlan1
```

```
Vxlan1 is up, line protocol is up (connected)
```

```
Hardware is Vxlan
```

```
Source interface is Loopback0 and is active with 3.3.3.3
```

```
Replication/Flood Mode is headend with Flood List Source: VCS
```

```
Remote MAC learning via VCS
```

```
Static vlan to vni mapping is
```

```
[100, 10000]
```

```
Dynamic vlan to vni mapping for 'vcs' is
```

```
[1120, 11120]
```

```
[1190, 11190]
```

```
[1174, 11174]
```

```
[1143, 11143]
```

```
[1145, 11145]
```

```
Headend replication flood vtep list is:
```

```
1120 3.3.3.3
```

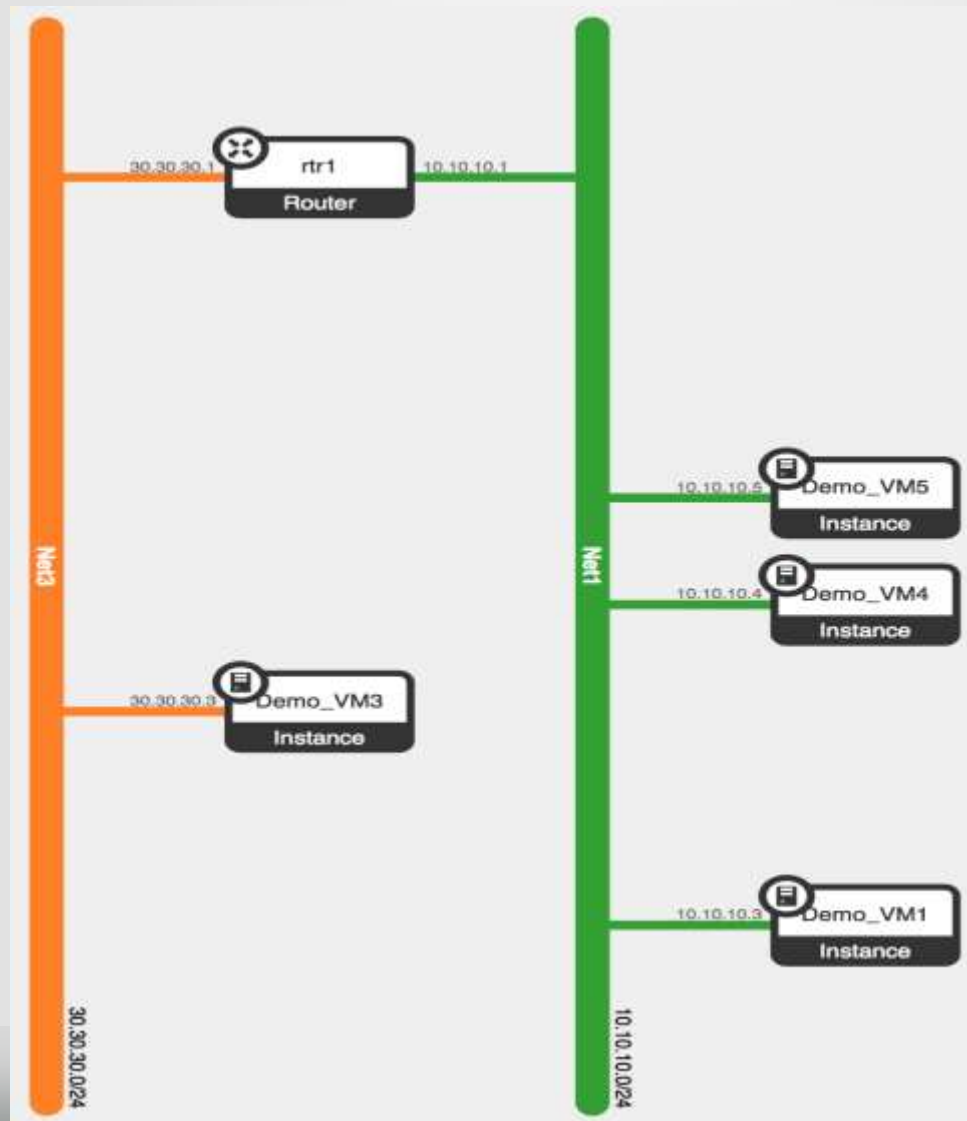
```
1143 3.3.3.3
```

```
1145 3.3.3.3      4.4.4.4
```

```
1174 3.3.3.3      4.4.4.4
```

```
1190 3.3.3.3      4.4.4.4
```

# L3 plugin - Arista is L3 router





# Create Network Net3

## Create Network

Network

Subnet \*

Subnet Details

Network Name

Net3

Admin State ?

UP

Create a new network. In addition, a subnet associated with the network can be created in the next panel.

Cancel

Back

Next >

## Network Details: Net3

### Network Overview

Name	Net3
ID	131b0b55-f4dc-4602-9510-6eb744a1bd3a
Project ID	c832cc8c4739426cbb0c08824fabe10b
Status	ACTIVE
Admin State	UP
Shared	No
External Network	No
MTU	Unknown
Provider Network	Network Type: vlan
	Physical Network: default
	Segmentation ID: 1143

# Create router which connect 2 networks

## Router Details

Set Gateway ▼

Overview

Interfaces

Static Routes

+ Add Interface

✕ Delete Interfaces

<input type="checkbox"/>	Name	Fixed IPs	Status	Type	Admin State	Actions
<input type="checkbox"/>	(2e2d1f16-f9fe)	30.30.30.1	Active	Internal Interface	UP	Delete Interface
<input type="checkbox"/>	(88540a4a-3542)	10.10.10.1	Active	Internal Interface	UP	Delete Interface


Displaying 2 items

## L3 Gateway on Arista switch

```
bleaf1#show run int vlan 1174
interface Vlan1174
  ip address 10.10.10.254/24
  ip virtual-router address 10.10.10.1
bleaf1#
bleaf1#
bleaf1#
bleaf1#show run int vlan 1143
interface Vlan1143
  ip address 30.30.30.254/24
  ip virtual-router address 30.30.30.1
```

```
bleaf1#show ip int bri | include 1174|1143
```

Vlan1143	30.30.30.254/24	up	up	1500
Vlan1174	10.10.10.254/24	up	up	1500



Thank-You

**ARISTA**